



哈爾濱工業大學
HARBIN INSTITUTE OF TECHNOLOGY



操作系统

Operating Systems

刘川意 副教授

哈尔滨工业大学（深圳）

2019年10月



Module 3: 同步 (Synchronization)

1. 信号的介绍 (Introduction to Signals)

CSAPP3e: 8.5

2. 同步 (Synchronization)

CSAPP3e: 12.4-12.7



Shared Variables in Threaded C Programs

- Question: Which variables in a threaded C program are shared?
 - The answer is not as simple as “*global variables are shared*” and “*stack variables are private*”
- *Definition:* A variable x is *shared* if and only if multiple threads reference some instance of x .
- Requires answers to the following questions:
 - What is the memory model for threads?
 - How are instances of variables mapped to memory?
 - How many threads might reference each of these instances?

Threads Memory Model

■ Conceptual model (概念模型):

- Multiple threads run within the context of a single process
- Each thread has its own separate thread context
 - ▶ Thread ID, stack, stack pointer, PC, condition codes, and GP registers
- All threads share the remaining process context
 - ▶ Code, data, heap, and shared library segments of the process virtual address space
 - ▶ Open files and installed handlers

■ Operationally, this model is not strictly enforced:

- Register (寄存器) values are truly separate and protected, but virtual memory is always shared
- Any thread can read and write the stack of any other thread

The mismatch between the conceptual (概念) and operation (操作) model is a source of confusion and errors

Example Program to Illustrate Sharing

```
6. char **ptr; /* global var */
```

```
33. int main()
```

```
34. {  
35.     long i;  
36.     pthread_t tid;  
37.     char *msgs[2] = {  
38.         "Hello from foo",  
39.         "Hello from bar"  
40.     };
```

```
41.     ptr = msgs;  
42.     for (i = 0; i < 2; i++)  
43.         Pthread_create(&tid,  
44.             NULL,  
45.             thread,  
46.             (void *)i);  
47.     Pthread_exit(NULL);  
48. }
```

sharing.c

```
7. void *thread(void *vargp)
```

```
8. {
```

```
9.     long myid = (long)vargp;
```

```
10.    static int cnt = 0;
```

```
11.    printf("[%ld]:  %s (cnt=%d)\n",  
12.        myid, ptr[myid], ++cnt);
```

```
13.    return NULL;  
14. }
```

Peer threads reference main thread's stack indirectly through global ptr variable

Mapping Variable Instances to Memory

■ Global variables

- *Def*: Variable declared outside of a function
- **Virtual memory contains exactly one instance of any global variable**

■ Local variables

- *Def*: Variable declared inside function without `static` attribute
- **Each thread stack contains one instance of each local variable**

■ Local static variables

- *Def*: Variable declared inside function with the `static` attribute
- **Virtual memory contains exactly one instance of any local static variable.**

Mapping Variable Instances to Memory

Global var: 1 instance (`ptr` [data])

```
6. char **ptr; /* global var */  
  
33. int main()  
34. {  
35.     long i;  
36.     pthread_t tid;  
37.     char *msgs[2] = {  
38.         "Hello from foo",  
39.         "Hello from bar"  
40.     };  
  
41.     ptr = msgs;  
42.     for (i = 0; i < 2; i++)  
43.         Pthread_create(&tid,  
44.             NULL,  
45.             thread,  
46.             (void *)i);  
47.     Pthread_exit(NULL);  
48. }
```

sharing.c

Local vars: 1 instance (`i.m`, `msgs.m`)

Local var: 2 instances (
 `myid.p0` [peer thread 0's stack],
 `myid.p1` [peer thread 1's stack]
)

```
7. void *thread(void *vargp)  
8. {  
9.     long myid = (long)vargp;  
10.    static int cnt = 0;  
  
11.    printf("[%ld]:  %s (cnt=%d)\n",  
12.        myid, ptr[myid], ++cnt);  
13.    return NULL;  
14. }
```

Local static var: 1 instance (`cnt` [data])

Shared Variable Analysis

■ Which variables are shared?

| <i>Variable instance</i> | <i>Referenced by main thread?</i> | <i>Referenced by peer thread 0?</i> | <i>Referenced by peer thread 1?</i> |
|--------------------------|-----------------------------------|-------------------------------------|-------------------------------------|
| <code>ptr</code> | yes | yes | yes |
| <code>cnt</code> | no | yes | yes |
| <code>i.m</code> | yes | no | no |
| <code>msgs.m</code> | yes | yes | yes |
| <code>myid.p0</code> | no | yes | no |
| <code>myid.p1</code> | no | no | yes |

■ Answer: A variable x is shared iff multiple threads reference at least one instance of x . Thus:

- `ptr`, `cnt`, and `msgs` are shared
- `i` and `myid` are **not** shared

badcnt.c: Improper Synchronization

```
7.  /* Global shared variable */
8.  volatile long cnt = 0; /* Counter */

17. int main(int argc, char **argv)
18. {
19.     long niters;
20.     pthread_t tid1, tid2;

21.     niters = atoi(argv[1]);
22.     Pthread_create(&tid1, NULL,
23.                   thread, &niters);
24.     Pthread_create(&tid2, NULL,
25.                   thread, &niters);
26.     Pthread_join(tid1, NULL);
27.     Pthread_join(tid2, NULL);

28.     /* Check result */
29.     if (cnt != (2 * niters))
30.         printf("BOOM! cnt=%ld\n",
31.               cnt);
32.     else
33.         printf("OK cnt=%ld\n", cnt);
34.     exit(0);
35. }
```

badcnt.c

```
38. /* Thread routine */
39. void *thread(void *vargp)
40. {
41.     long i, niters =
42.         *((long *)vargp);
43.
44.     for (i = 0; i < niters; i++)
45.         cnt++;
46.
47.     return NULL;
48. }
```

```
[zs_cao@localhost conc]$ ./badcnt 10000
OK cnt=20000
[zs_cao@localhost conc]$ ./badcnt 10000
BOOM! cnt=17302
[zs_cao@localhost conc]$ ./badcnt 10000
OK cnt=20000
```

线程并发执行的问题



Assembly Code for Counter Loop

■ 编译: 来自CSAPP第7章链接

- gcc -s badcnt.c -o badcnt.s
- vim badcnt.s

```
for (i = 0; i < niters; i++)  
    cnt++;
```

```
94      movq    %rdi, -24(%rbp)  
95      movq    -24(%rbp), %rax  
96      movq    (%rax), %rax  
97      movq    %rax, -8(%rbp)  
98      movq    $0, -16(%rbp)  
99      jmp     .L6
```

100 .L7:

```
101      movq    cnt(%rip), %rax  
102      addq    $1, %rax  
103      movq    %rax, cnt(%rip)  
104      addq    $1, -16(%rbp)
```

105 .L6:

```
106      movq    -16(%rbp), %rax  
107      cmpq    -8(%rbp), %rax  
108      jl      .L7  
109      movl    $0, %eax  
110      popq    %rbp
```

H_i : Head

L_i : Load cnt

U_i : Update cnt

S_i : Store cnt

T_i : Tail

Assembly Code for Counter Loop

■ 汇编：来自CSAPP第7章 链接

- gcc -c badcnt.s -o badcnt.o
- objdump -dx badcnt.o

```
for (i = 0; i < niters; i++)
    cnt++;
```

000000000000000ed <thread>:

| | | | | | |
|---------------|----------------------|------|-------------------|---------|--|
| ed: | 55 | push | %rbp | | |
| ee: | 48 89 e5 | mov | %rsp,%rbp | | |
| f1: | 48 89 7d e8 | mov | %rdi,-0x18(%rbp) | | |
| f5: | 48 8b 45 e8 | mov | -0x18(%rbp),%rax | | |
| f9: | 48 8b 00 | mov | (%rax),%rax | | |
| fc: | 48 89 45 f8 | mov | %rax,-0x8(%rbp) | | |
| 100: | 48 c7 45 f0 00 00 00 | movq | \$0x0,-0x10(%rbp) | | |
| 107: | 00 | | | | |
| 108: | eb 17 | jmp | 121 <thread+0x34> | | |
| 10a: | 48 8b 05 00 00 00 00 | mov | 0x0(%rip),%rax | # 111 | |
| <thread+0x24> | | | | | |
| | | 10d: | R_X86_64_PC32 | cnt-0x4 | |
| 111: | 48 83 c0 01 | add | \$0x1,%rax | | |
| 115: | 48 89 05 00 00 00 00 | mov | %rax,0x0(%rip) | # 11c | |
| <thread+0x2f> | | | | | |
| | | 118: | R_X86_64_PC32 | cnt-0x4 | |
| 11c: | 48 83 45 f0 01 | addq | \$0x1,-0x10(%rbp) | | |
| 121: | 48 8b 45 f0 | mov | -0x10(%rbp),%rax | | |
| 125: | 48 3b 45 f8 | cmp | -0x8(%rbp),%rax | | |
| 129: | 7c df | j1 | 10a <thread+0x1d> | | |
| 12b: | b8 00 00 00 00 | mov | \$0x0,%eax | | |
| 130: | 5d | pop | %rbp | | |
| 131: | c3 | retq | | | |



Assembly Code for Counter Loop

■ 链接: 来自CSAPP第7章 链接

- gcc -o badcnt.c -o badcnt -lpthread
- objdump -d badcnt

```
for (i = 0; i < niters; i++)
    cnt++;
```

00000000000000957 <thread>:

| | | | | |
|------|----------------------|------|---------------------|--|
| 957: | 55 | push | %rbp | |
| 958: | 48 89 e5 | mov | %rsp,%rbp | |
| 95b: | 48 89 7d e8 | mov | %rdi,-0x18(%rbp) | |
| 95f: | 48 8b 45 e8 | mov | -0x18(%rbp),%rax | |
| 963: | 48 8b 00 | mov | (%rax),%rax | |
| 966: | 48 89 45 f8 | mov | %rax,-0x8(%rbp) | |
| 96a: | 48 c7 45 f0 00 00 00 | movq | \$0x0,-0x10(%rbp) | |
| 971: | 00 | | | |
| 972: | eb 17 | jmp | 98b <thread+0x34> | |
| 974: | 48 8b 05 b5 06 20 00 | mov | 0x2006b5(%rip),%rax | |
| | # 201030 <cnt> | | | |
| 97b: | 48 83 c0 01 | add | \$0x1,%rax | |
| 97f: | 48 89 05 aa 06 20 00 | mov | %rax,0x2006aa(%rip) | |
| | # 201030 <cnt> | | | |
| 986: | 48 83 45 f0 01 | addq | \$0x1,-0x10(%rbp) | |
| 98b: | 48 8b 45 f0 | mov | -0x10(%rbp),%rax | |
| 98f: | 48 3b 45 f8 | cmp | -0x8(%rbp),%rax | |
| 993: | 7c df | jl | 974 <thread+0x1d> | |
| 995: | b8 00 00 00 00 | mov | \$0x0,%eax | |
| 99a: | 5d | pop | %rbp | |
| 99b: | c3 | retq | | |

怎么计算?

H_i

L_i

U_i

S_i


T_i

Concurrent Execution(并发执行)


■ *Key idea:* In general, any sequentially consistent interleaving is possible, but some give an unexpected result!

- I_i denotes that thread i executes instruction I
- $\%rdx_i$ is the content of $\%rdx$ in thread i 's context

| i (thread) | $instr_i$ | $\%rdx_1$ | $\%rdx_2$ | cnt |
|--------------|-----------|-----------|-----------|-----|
| 1 | H_1 | - | - | 0 |
| 1 | L_1 | 0 | - | 0 |
| 1 | U_1 | 1 | - | 0 |
| 1 | S_1 | 1 | - | 1 |
| 2 | H_2 | - | - | 1 |
| 2 | L_2 | - | 1 | 1 |
| 2 | U_2 | - | 2 | 1 |
| 2 | S_2 | - | 2 | 2 |
| 2 | T_2 | - | 2 | 2 |
| 1 | T_1 | 1 | - | 2 |



Thread 1
critical section



Thread 2
critical section

OK

Concurrent Execution (cont)

- Incorrect ordering: two threads increment the counter, but the result is 1 instead of 2


| i (thread) | instr _i | %rdx ₁ | %rdx ₂ | cnt |
|------------|--------------------|-------------------|-------------------|-----|
| 1 | H ₁ | - | - | 0 |
| 1 | L ₁ | 0 | - | 0 |
| 1 | U ₁ | 1 | - | 0 |
| 2 | H ₂ | - | - | 0 |
| 2 | L ₂ | - | 0 | 0 |
| 1 | S ₁ | 1 | - | 1 |
| 1 | T ₁ | 1 | - | 1 |
| 2 | U ₂ | - | 1 | 1 |
| 2 | S ₂ | - | 1 | 1 |
| 2 | T ₂ | - | 1 | 1 |

S1应该在L2之前执行

Oops!

Progress Graphs (进度图)

Thread 2



A *progress graph* depicts the discrete *execution state space* of concurrent threads.


Each axis corresponds to the sequential order of instructions in a thread.

Each point corresponds to a possible *execution state* ($\text{Inst}_1, \text{Inst}_2$).

E.g., (L_1, S_2) denotes state where thread 1 has completed L_1 and thread 2 has completed S_2 .

Trajectories in Progress Graphs

Thread 2




A *trajectory* (轨道) is a sequence of legal state transitions (转换) that describes one possible concurrent execution of the threads.

Example:

$H_1, L_1, U_1, H_2, L_2, S_1, T_1, U_2, S_2, T_2$

Critical Sections and Unsafe Regions




L, U, and S form a *critical section* (临界区) with respect to the shared variable `cnt`

Instructions in critical sections (wrt some shared variable) should not be interleaved

Sets of states where such interleaving occurs form *unsafe regions*

Critical Sections and Unsafe Regions



Def: A trajectory (轨道) is *safe* iff it does not enter any unsafe region

Claim: A trajectory is correct (wrt cnt) iff it is safe

Enforcing Mutual Exclusion

- *Question:* How can we guarantee a safe trajectory?
- Answer: We must **synchronize** (同步) the execution of the threads so that they can never have an unsafe trajectory.
 - i.e., need to guarantee **mutually exclusive access** (互斥地访问) for each critical section.
- Classic solution:
 - Semaphores (信号量) (Edsger Dijkstra)
- Other approaches (out of our scope)
 - Mutex and condition variables (Pthreads)
 - Monitors (Java)

Semaphores (信号量)

- **Semaphore:** non-negative global integer synchronization variable. Manipulated by P and V operations. **Semaphore invariant:** ($s \geq 0$)
- $P(s)$
 - If s is nonzero, then decrement s by 1 and return immediately.
 - ▶ Test and decrement operations occur atomically (indivisibly)
 - If s is zero, then suspend thread until s becomes nonzero and the thread is restarted by a V operation.
 - After restarting, the P operation decrements s and returns control to the caller.
- $V(s)$:
 - Increment s by 1.
 - ▶ Increment operation occurs atomically
 - If there are any threads blocked in a P operation waiting for s to become non-zero, then restart exactly one of those threads, which then completes its P operation by decrementing s .

```
1. P(s) :  
2.     if s>0 s--;  
3.     if s==0 do block;  
4. V(s) :  
5.     s++;  
6.     if any threads blocked  
   by P operation then wakeup
```

C Semaphore Operations

Pthreads functions:

```
#include <semaphore.h>

int sem_init(sem_t *s, int pshared, unsigned int val);}
/* s为指向信号量结构的一个指针;
   pshared不为0时此信号量在进程间共享, 否则只能为当前进程的所有线程
   共享;
   val给出了信号量的初始值。
   成功返回0, 错误返回-1
*/
int sem_wait(sem_t *s); /* P(s) */
int sem_post(sem_t *s); /* V(s) */
```

CS:APP wrapper functions:

```
94. void P(sem_t *sem)
95. {
96.     if (sem_wait(sem) < 0)
97.         unix_error("P error");
98. }
```

goodcnt.c

```
100. void V(sem_t *sem)
101. {
102.     if (sem_post(sem) < 0)
103.         unix_error("V error");
104. }
```

goodcnt.c

badcnt.c: Improper Synchronization

```
7.  /* Global shared variable */
8.  volatile long cnt = 0; /* Counter */

17. int main(int argc, char **argv)
18. {
19.     long niters;
20.     pthread_t tid1, tid2;

21.     niters = atoi(argv[1]);
22.     Pthread_create(&tid1, NULL,
23.         thread, &niters);
24.     Pthread_create(&tid2, NULL,
25.         thread, &niters);
26.     Pthread_join(tid1, NULL);
27.     Pthread_join(tid2, NULL);

28.     /* Check result */
29.     if (cnt != (2 * niters))
30.         printf("BOOM! cnt=%ld\n",
31.             cnt);
32.     else
33.         printf("OK cnt=%ld\n", cnt);
34.     exit(0);
35. }
```

badcnt.c

```
38. /* Thread routine */
39. void *thread(void *vargp)
40. {
41.     long i, niters =
42.         *((long *)vargp);
43.
44.     for (i = 0; i < niters; i++)
45.         cnt++;
46.
47.     return NULL;
48. }
```

How can we fix this using semaphores?

Using Semaphores for Mutual Exclusion

■ Basic idea:

- Associate a unique semaphore *mutex*, initially 1, with each shared variable (or related set of shared variables).
- Surround corresponding critical sections (临界区) with $P(mutex)$ and $V(mutex)$ operations.

■ Terminology:

- *Binary semaphore (二元信号量)*: semaphore whose value is always 0 or 1
- *Mutex (互斥)*: binary semaphore used for mutual exclusion
 - ▶ P operation: “locking” the mutex
 - ▶ V operation: “unlocking” or “releasing” the mutex
 - ▶ “Holding” a mutex: locked and not yet unlocked.
- *Counting semaphore (计数信号量)*: used as a counter for set of available resources.

goodcnt.c: Proper Synchronization

- Define and initialize a mutex for the shared variable `cnt`:

```
volatile long cnt = 0; /* Counter */
sem_t mutex;          /* Semaphore that protects cnt */

Sem_init(&mutex, 0, 1); /* mutex = 1 */
```

- **Surround** critical section with *P* and *V*:

```
for (i = 0; i < niters; i++) {
    P(&mutex);
    cnt++;
    V(&mutex);
}
```

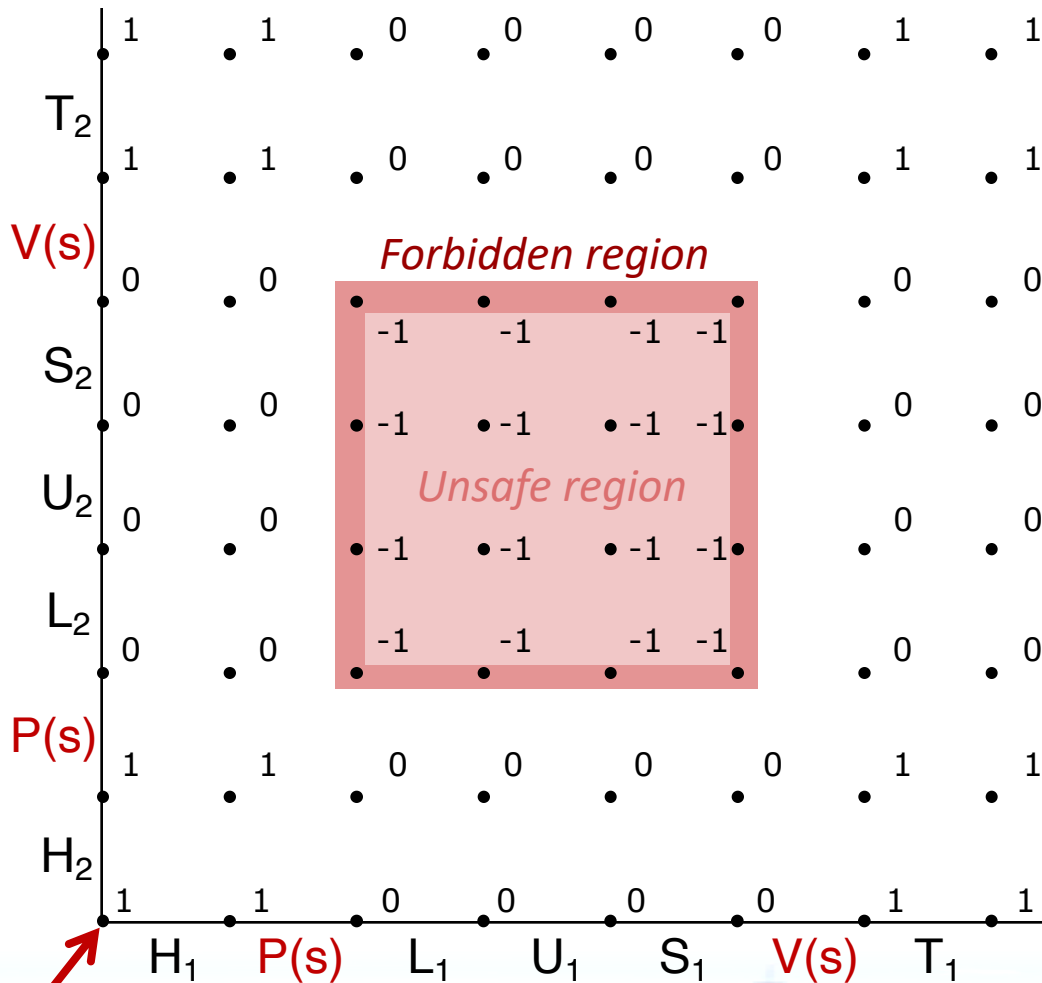
goodcnt.c

```
linux> ./goodcnt 10000
OK cnt=20000
linux> ./goodcnt 10000
OK cnt=20000
linux>
```

Warning: It's orders of magnitude slower than
`badcnt.c`.

Why Mutexes Work

Thread 2



Provide mutually exclusive access to shared variable by surrounding critical section with P and V operations on semaphore s (initially set to 1)

Semaphore invariant creates a *forbidden region* (禁止区域) that encloses unsafe region and that cannot be entered by any trajectory.

Initially
 $s = 1$

Thread 1



Summary

- Programmers need a clear model of how variables are shared by threads.
- Variables shared by multiple threads must be protected to ensure mutually exclusive access.
- Semaphores are a fundamental mechanism for enforcing mutual exclusion.
- 使用信号量的函数执行时间较未使用信号量的时间要长
 - 以goodcnt和badcnt为例, 参数为1000000时, goodcnt运行时长为badcnt时长的17倍

```
[zs_cao@localhost conc]$ time ./badcnt 1000000
OK cnt=2000000


real    0m0.013s
user    0m0.004s
sys     0m0.008s
[zs_cao@localhost conc]$ time ./goodcnt 1000000
OK cnt=2000000

real    0m0.225s
user    0m0.187s
sys     0m0.145s
[zs_cao@localhost conc]$
```

Using Semaphores to Coordinate Access to Shared Resources

- Basic idea: Thread uses a semaphore operation to notify another thread that some condition has become true
 - Use counting semaphores to keep track of resource state and to notify other threads
 - Use mutex to protect access to resource
- Two classic examples:
 - The Producer-Consumer Problem
 - The Readers-Writers Problem

Producer-Consumer Problem



■ Common synchronization pattern:

- Producer waits for empty **slot**, inserts item in buffer, and notifies consumer
- Consumer waits for **item**, removes it from buffer, and notifies producer

■ Examples

- Multimedia processing:
 - ▶ Producer creates MPEG video frames, consumer renders them
- Event-driven graphical user interfaces
 - ▶ Producer detects mouse clicks, mouse movements, and keyboard hits and inserts corresponding events in buffer
 - ▶ Consumer retrieves events from buffer and paints the display

Producer-Consumer on an n -element Buffer

- Requires a mutex and two counting semaphores:
 - `mutex`: enforces mutually exclusive access to the the buffer
 - `slots`: counts the available slots in the buffer
 - `items`: counts the available items in the buffer
- Implemented using a shared buffer package called `sbuf`.

sbuf Package - Declarations

```
#include "csapp.h"

typedef struct {
    int *buf;           /* Buffer array */
    int n;              /* Maximum number of slots */
    int front;          /* buf[(front+1)%n] is first item */
    int rear;           /* buf[rear%n] is last item */
    sem_t mutex;        /* Protects accesses to buf */
    sem_t slots;        /* Counts available slots */
    sem_t items;        /* Counts available items */
} sbuf_t;

void sbuf_init(sbuf_t *sp, int n);      /* 初始化信号量 */
void sbuf_deinit(sbuf_t *sp);          /* 释放buf */
void sbuf_insert(sbuf_t *sp, int item); /* 生产者线程 */
int sbuf_remove(sbuf_t *sp);            /* 消费者线程 */
```

sbuf.h

sbuf Package - Implementation

Initializing and deinitializing a shared buffer:

```
/* Create an empty, bounded, shared FIFO buffer with n slots */
void sbuf_init(sbuf_t *sp, int n)
{
    sp->buf = Calloc(n, sizeof(int));
    sp->n = n;                                /* Buffer holds max of n items */
    sp->front = sp->rear = 0;                 /* Empty buffer iff front == rear */
}

/* Binary semaphore for locking */
Sem_init(&sp->mutex, 0, 1);

/* Initially, buf has n empty slots */
Sem_init(&sp->slots, 0, n);

/* Initially, buf has 0 items */
Sem_init(&sp->items, 0, 0);
}

/* Clean up buffer sp */
void sbuf_deinit(sbuf_t *sp)
{
    Free(sp->buf);
}
```

sbuf.c

sbuf Package - Implementation

Inserting an item into a shared buffer:

```
/* Insert item onto the rear of shared buffer sp */
void sbuf_insert(sbuf_t *sp, int item)
{
    P(&sp->slots);                /* Wait for available slot
*/
    P(&sp->mutex);                /* Lock the buffer */
    sp->buf[(++sp->rear)%(sp->n)] = item; /* Insert the item */
    V(&sp->mutex);                /* Unlock the buffer */
    V(&sp->items);                /* Announce available item
*/
}
```

sbuf.c

sbuf Package - Implementation

Removing an item from a shared buffer:

```
/* Remove and return the first item from buffer sp */
int sbuf_remove(sbuf_t *sp)
{
    int item;
    P(&sp->items);                               /* Wait for available item */
    P(&sp->mutex);                                /* Lock the buffer */
    item = sp->buf[(++sp->front)%(sp->n)]; /* Remove the item */
    V(&sp->mutex);                                /* Unlock the buffer */
    V(&sp->slots);                                /* Announce available slot */
    return item;
}
```

sbuf.c